

A Parallel, Generic Sorting Library in Go

QOSORT

Shaojie Bai
Yutong Chen

QOSORT

- ▶ Portable library
- ▶ Supports generic data types
- ▶ Optimized for multi-core
- ▶ Scalable with resources

CHALLENGES

- ▶ Lack of explicit hardware scheduling
- ▶ Lack of parallel API support
- ▶ Enormity of data size
- ▶ Complications associated with generic data types

SORT IN GO


sort.Interface

```
type Interface interface {  
    // Len is the number of elements in the collection.  
    Len() int  
    // Less reports whether the element with  
    // index i should sort before the element with index j.  
    Less(i, j int) bool  
    // Swap swaps the elements with indexes i and j.  
    Swap(i, j int)  
}
```

QUICKSORT

- ▶ **Quick** recap
- ▶ Two factors at stake:
 - Pivot
 - Scheduling

QUICKSORT V1.0 – “NAIVE”




QUICKSORT TASK-BASED SCHEDULING


Splitting into 2 Parts

(i,j)


New Task with Indices (i,j)



Sample and Locate Pivot




Split Array into Two Parts Using Pivot




Iterate on Larger Part (i,m)

Send Smaller Part (n,j) to Task Queue

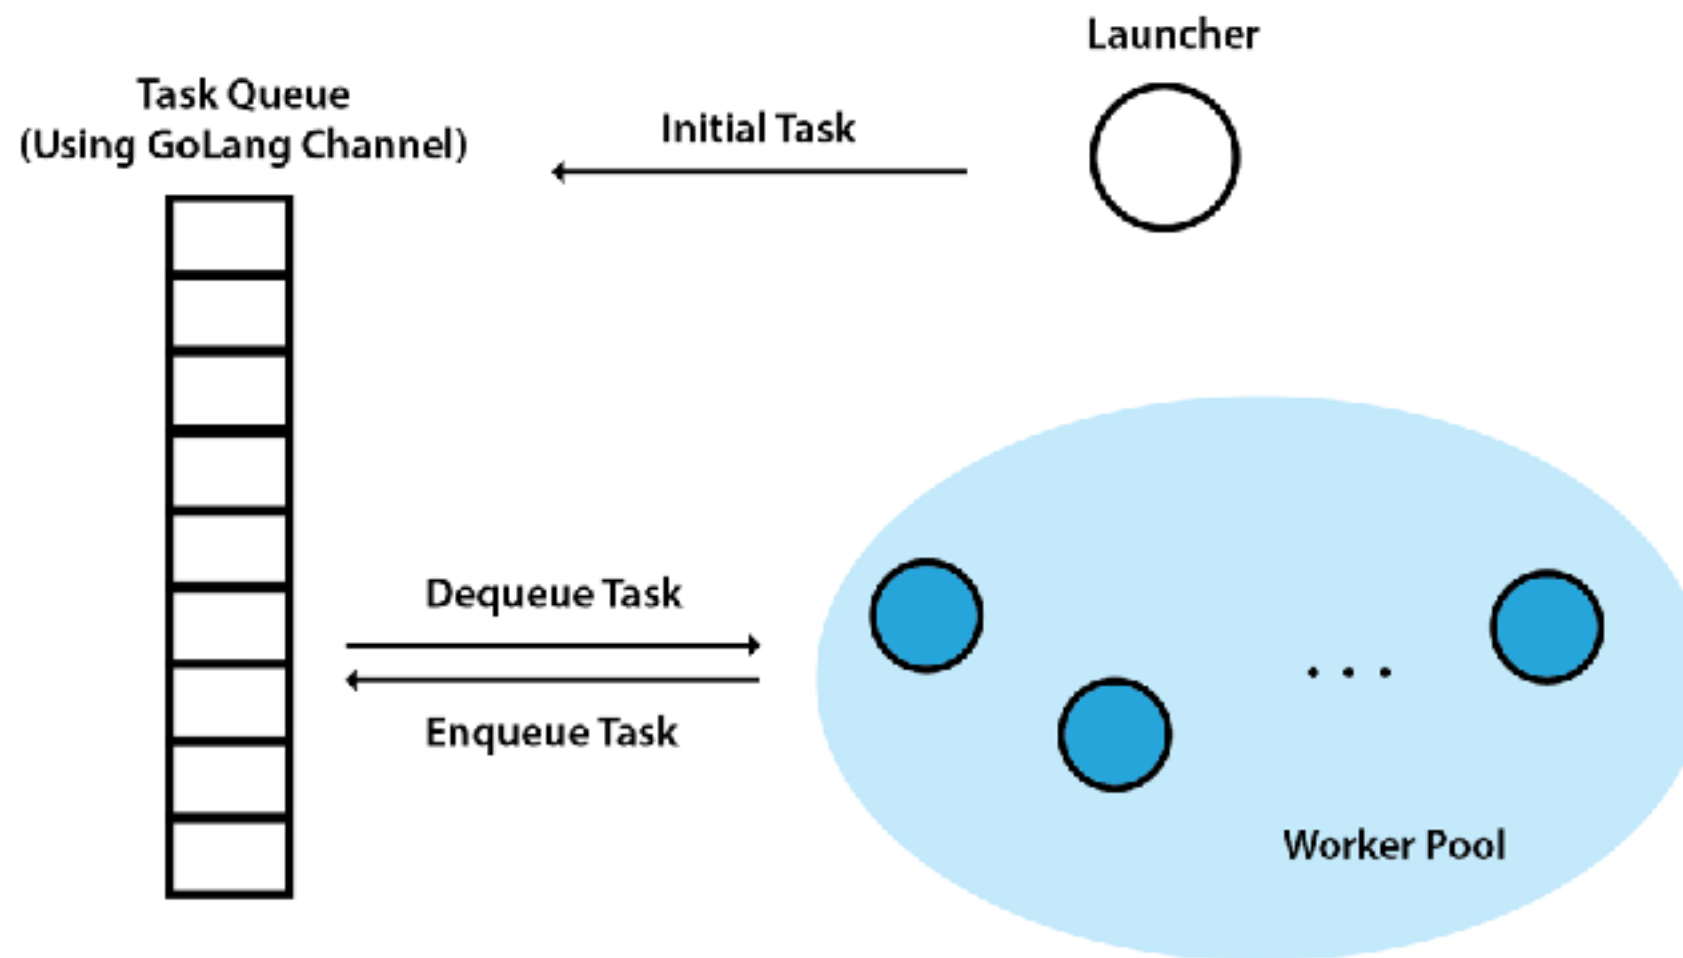


Task Queue

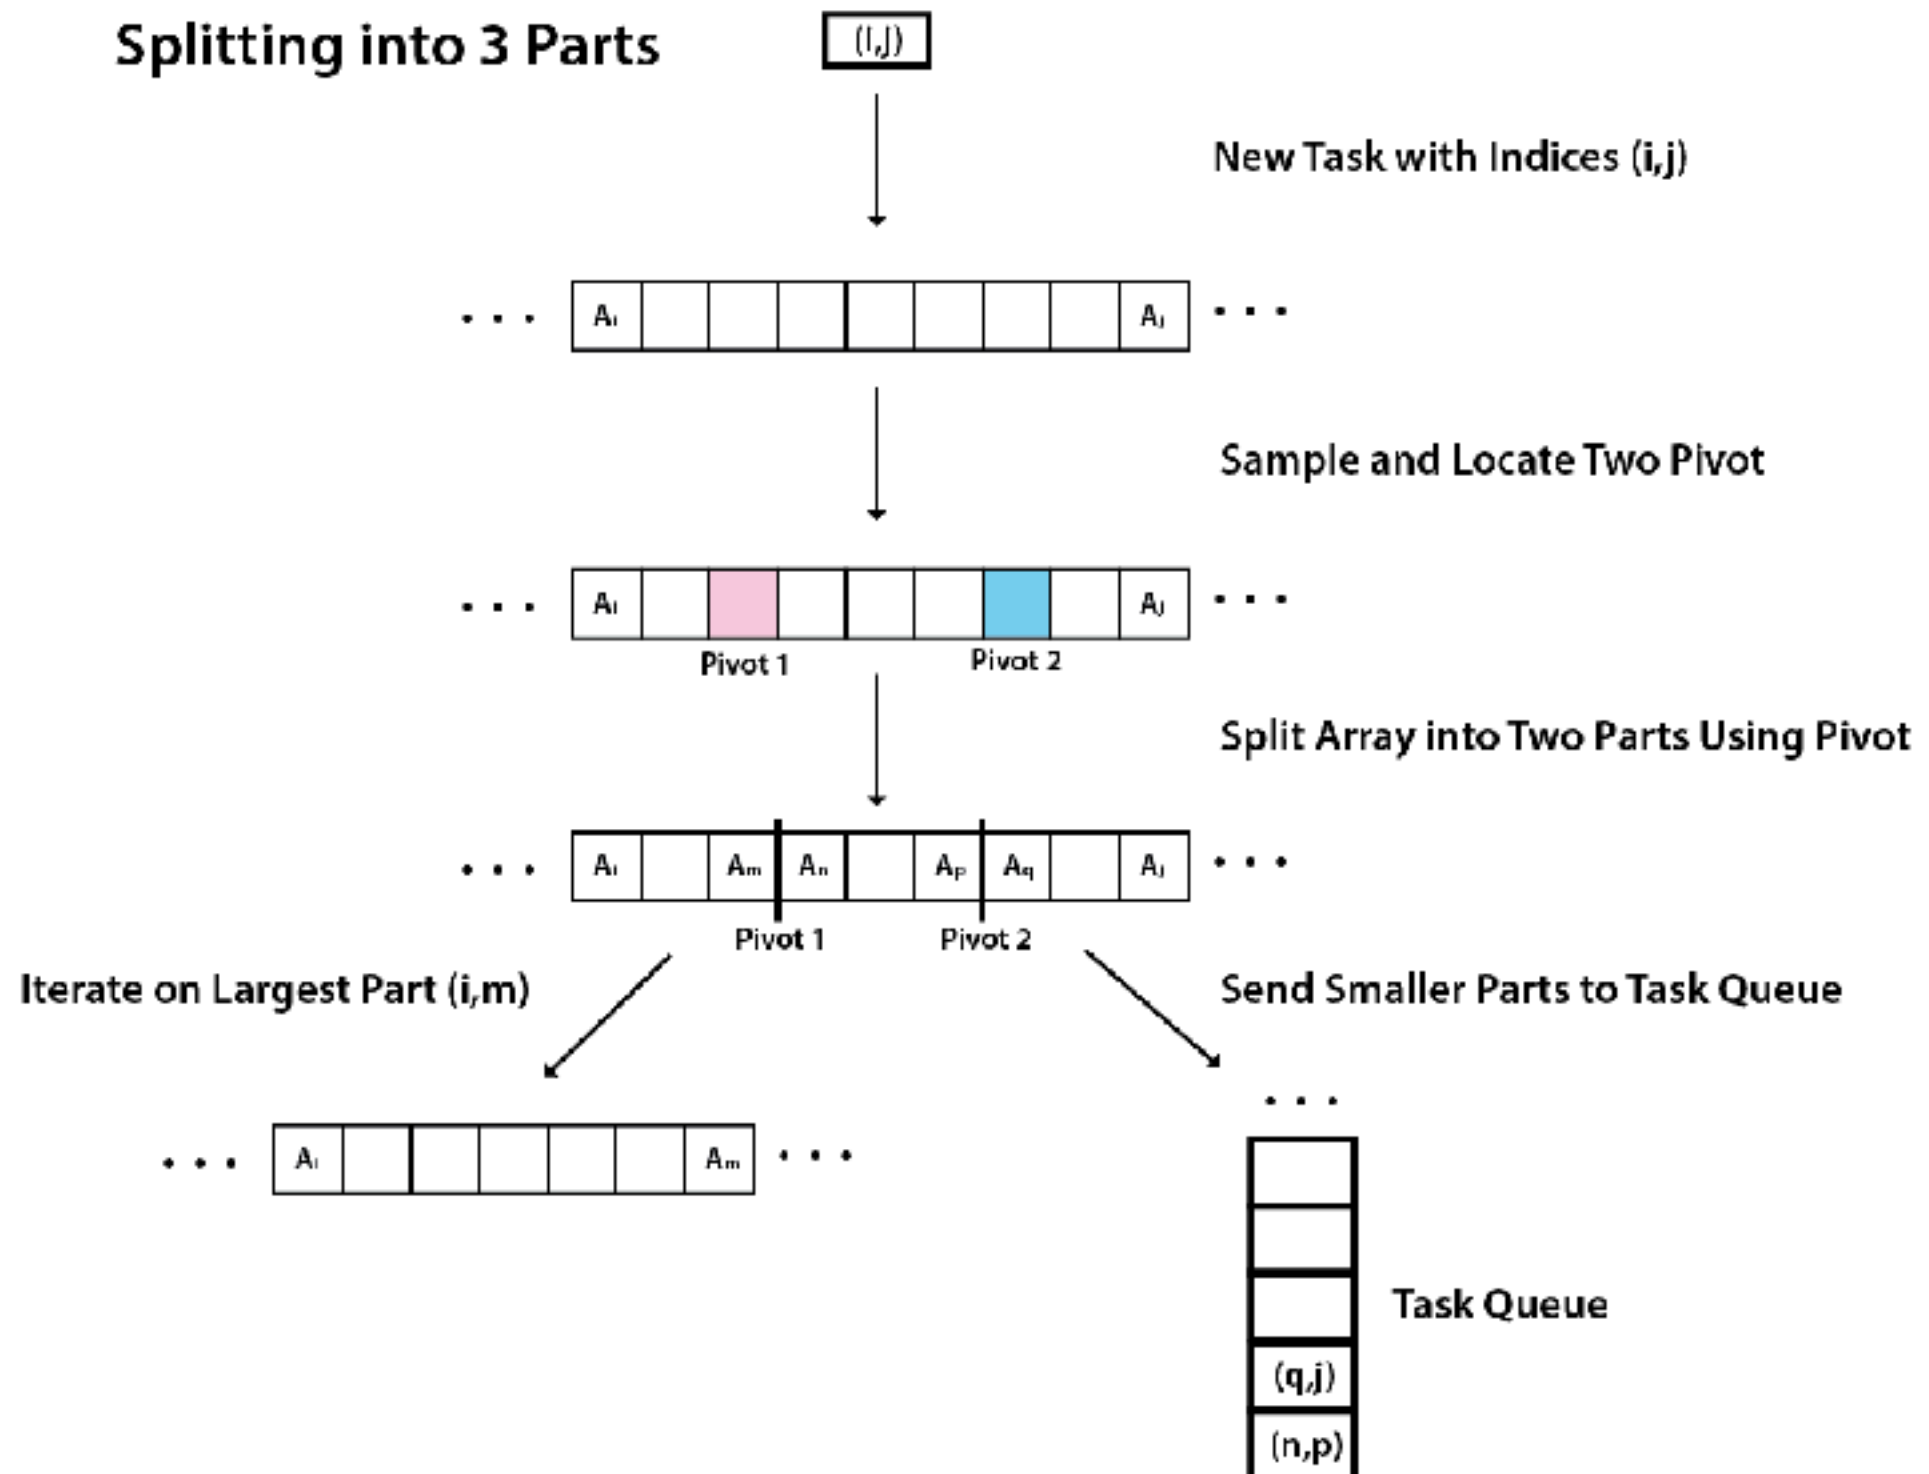
QUICKSORT V2.0 – “MASTER DISPATCH”



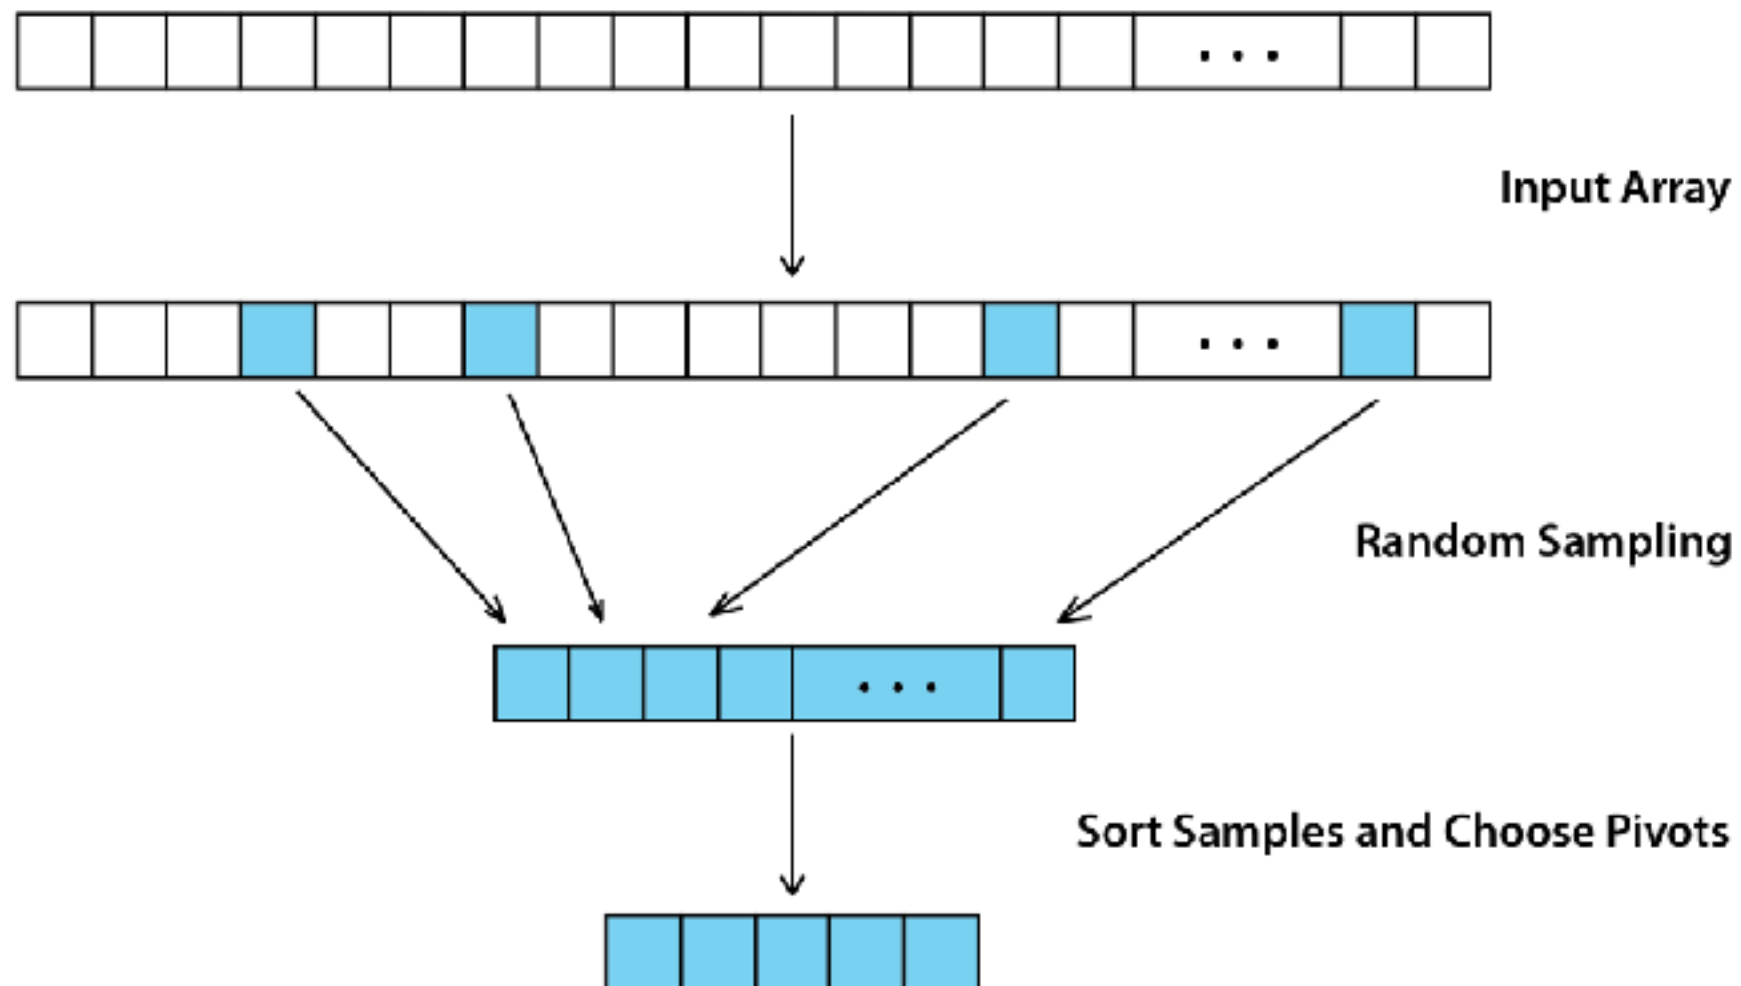
QUICKSORT V3.0 – “ONE QUEUE”



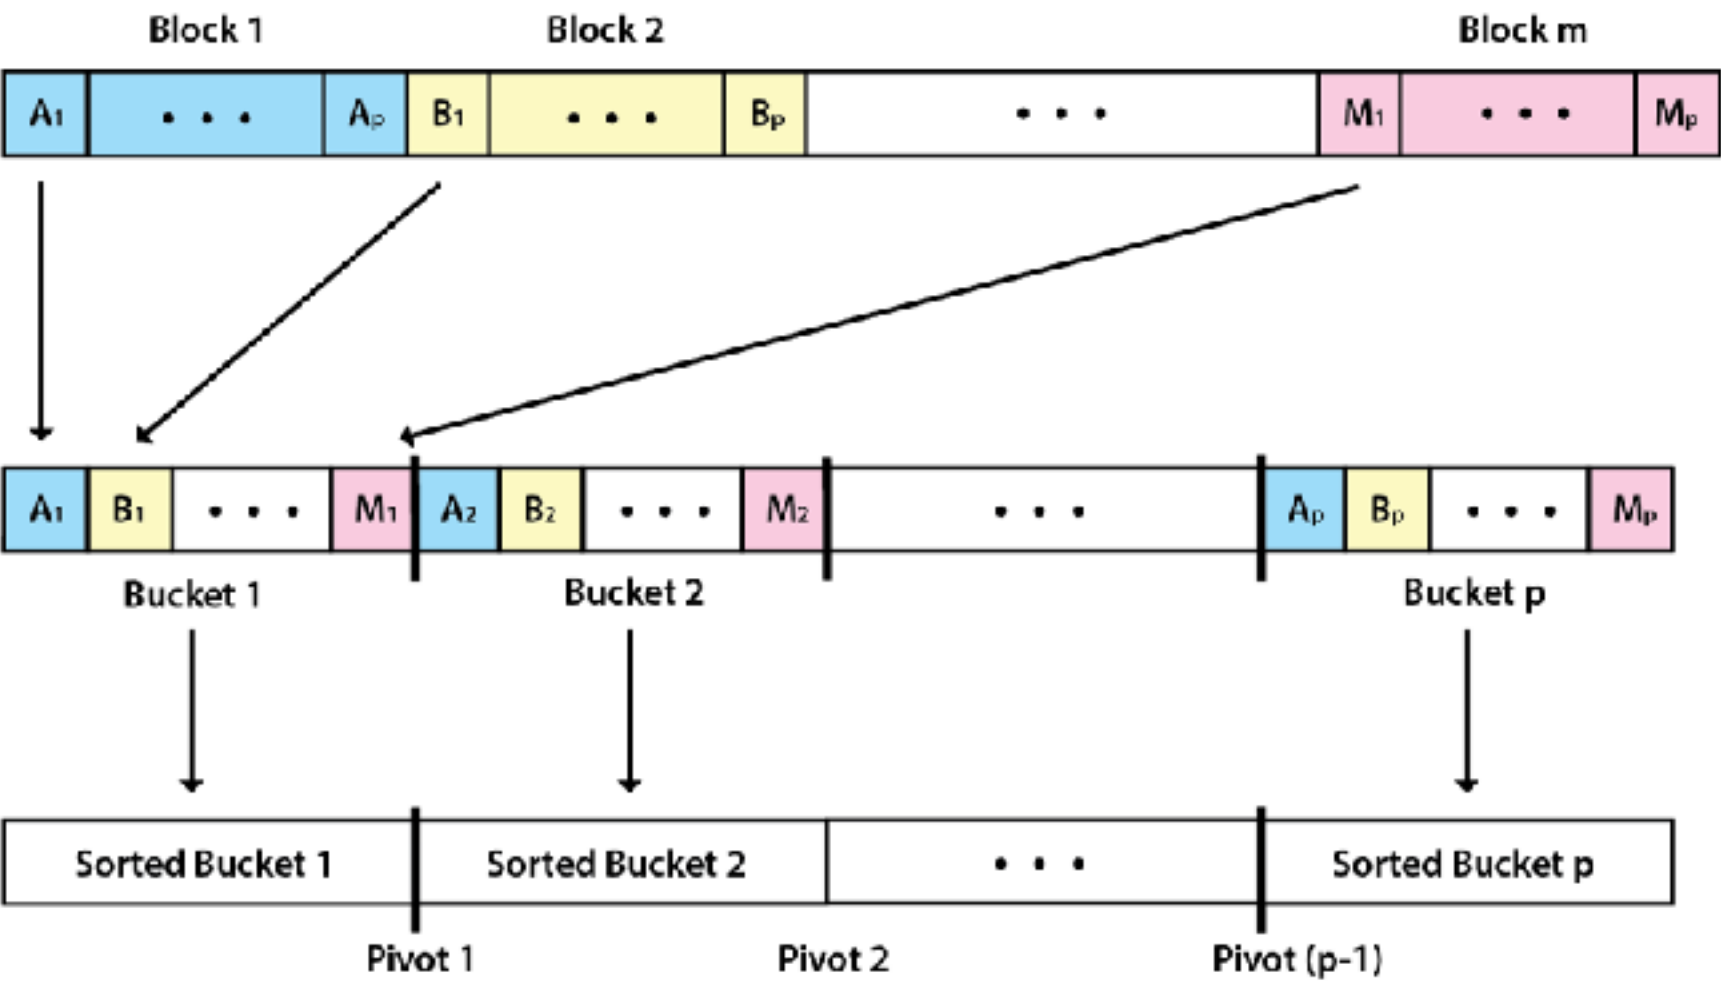
QUICKSORT V4.0 - "SPLIT-BY-3"



SAMPLESORT – SAMPLE



SAMPLESORT – BLOCK TRANSPOSE



Transpose Blocks to Buckets

Sort Each Bucket

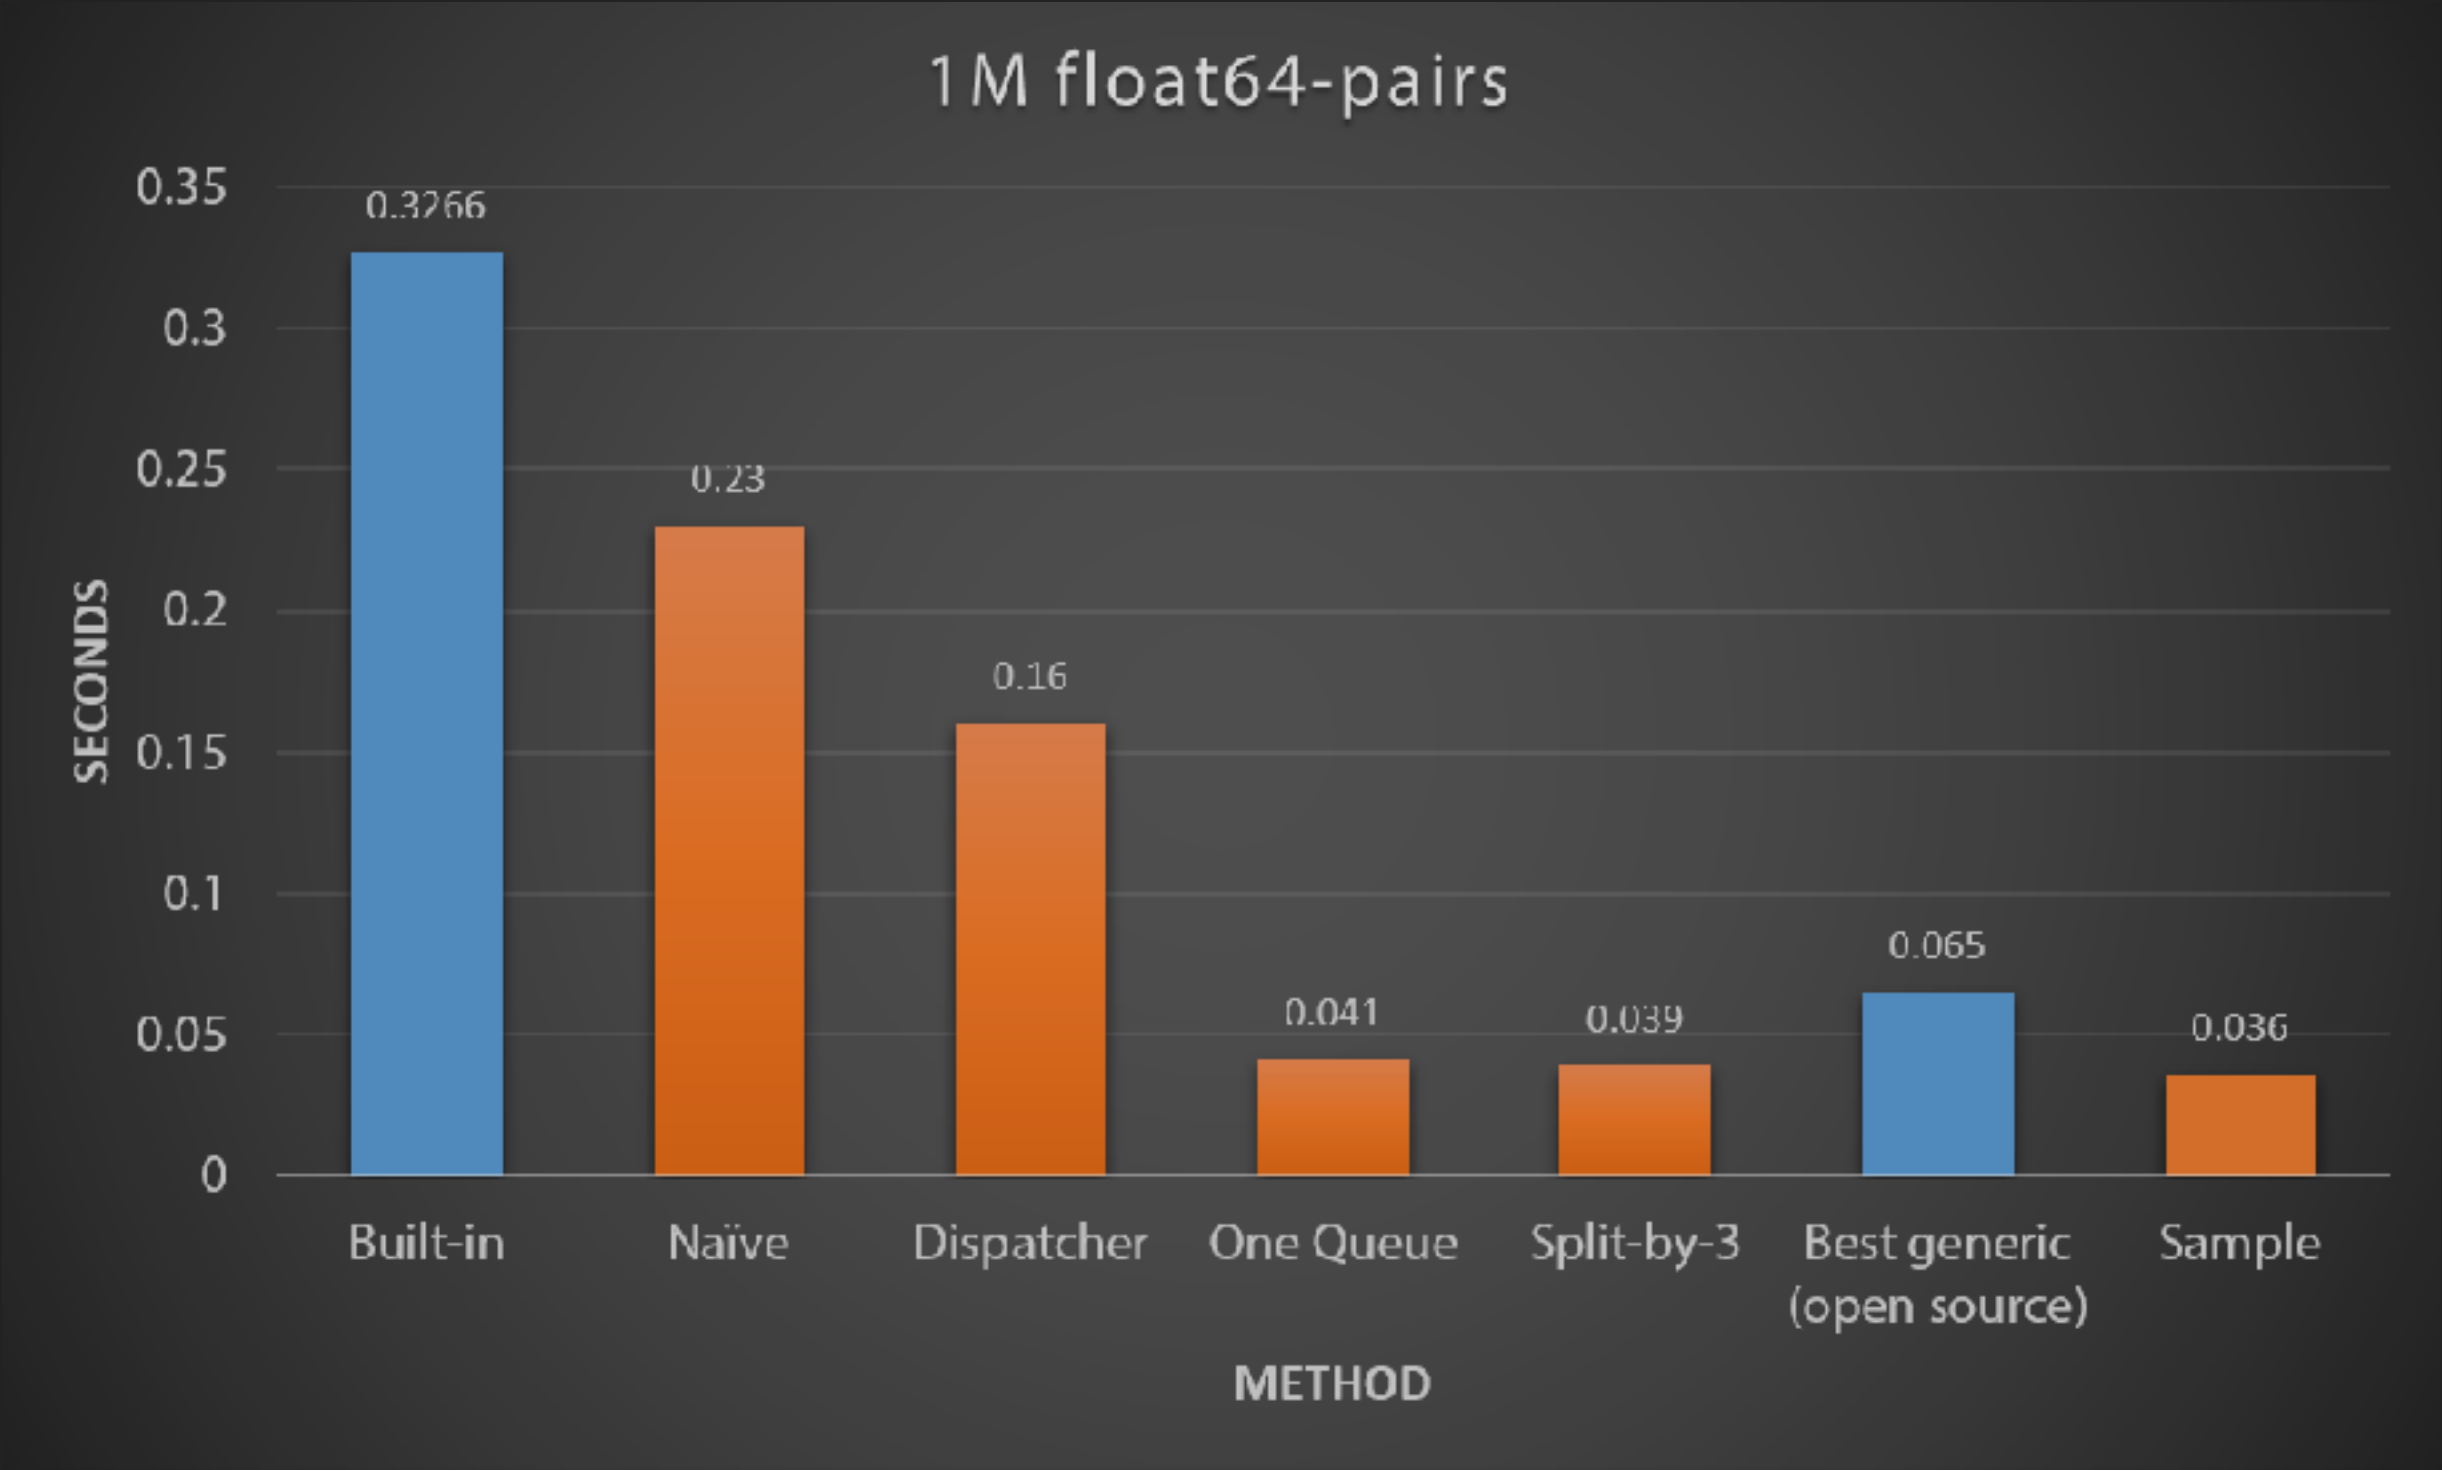
IMPROVEMENT OVER QUICKSORT

- ▶ Highly parallelizable
- ▶ Even better load balancing
- ▶ Parallelism starts early
- ▶ Spatial locality


TEST SETTING

- ▶ UNIX Machine - 2x Intel(R) Xeon(R) E5-2680 v2 (20 cores)
- ▶ Go 1.4.2 Linux/AMD64
- ▶ 100,000,000 float64 key-value pairs

PERFORMANCE



PERFORMANCE (CONT'D)



QUESTIONS?

- ▶ Questions for you consideration:
 - ▶ What's the size of worker pool ?
 - ▶ "Base case" for each algorithm ?
 - ▶ Why the name "QoSort" ?
 - ▶ Why is Kayvon replying my comment at 2am ?!
 - ▶ ...